

Using line features for 3D face registration

Fabian Brix

18.06.2013

Abstract

In this bachelor thesis we attempt to modify the existing face registration pipeline for the morphable face model of Prof. Thomas Vetter by using a registration algorithm developed by PD Marcel Lthi at the University of Basel. **ALTERNATIVE:** In this bachelor thesis we discuss the construction of a face registration pipeline using an algorithm based on a vector-valued gaussian process and at the same time attempting to ensure registration quality through the use of contours marking important parts of the face - referred to as line features.

The algorithm is capable of mapping any two shapes on to one another. All that is needed is a set of corresponding points on the two shapes. Different constraints to the displacement field can be applied through regularisation.

The aim of this bachelor thesis is more specifically to apply this general algorithm for point correspondences to scanned face data, that is to implement feasible registration of face scans onto the mean face of the morphable model. In order to achieve this we mark important parts of the face meshes not only with point landmarks, but also structures and organs (eyebrows, eyes, ears) with lines - line features - and thereby to create further correspondences for the algorithm to perform better by. Instead of using sparse points of key features points of the face we mark complex features, e.g. the eyes, with contour lines - line features in order to create further correspondences

These line features are marked by hand using bzier curves on three 2D images to the front, left and right of the 3D face. In order to utilize them, however, they have to be projected on to the computed mesh of the face that was recorded by a 3D scanner. These meshes have holes in the region of the eyes and the ears rendering the projected line features useless at first. This thesis first gives an overview over the morphable model and the face registration pipeline, then goes on to obtaining 3D points from the 2D line features, to explain the theory behind the general algorithm and in the main part discusses the problems and solutions we encountered trying to optimize the algorithm for and without line features for the face registration process.

Contents

39	Contents	2
40	1 Introduction	3
41	1.1 Problem Statement	3
42	1.2 Review Literature	3
43	2 3D Model Building	5
44	2.1 3D Morphable Model	5
45	2.2 Prerequisite Data	6
46	2.3 Finding Correspondences	6
47	3 Gaussian Processes in 3D Face Registration	7
48	3.1 Stochastic Processes	7
49	3.2 Gaussian Processes	7
50	3.3 Gaussian Process Regression	9
51	3.4 Application to 3D Face Meshs	10
52	3.5 Fitting & Optimization	12
53	3.6 Robust Loss Functions	13
54	4 Registration Pipeline using Line Features	15
55	4.1 Line Features	15
56	4.2 Sampling 3D Points from 2D Line Features	16
57	4.3 Preparing the Mean Mesh	21
58	4.4 Rigid Mesh Alignment	21
59	4.5 Prior Model	22
60	4.6 Posterior Model	22
61	4.7 Fitting	22
62	4.8 Optimizing the Loss Function	22
63	4.9 Varying the Variances	22

Chapter 1

Introduction

1.1 Problem Statement

So es bizzeli alles schriebe 1. Use Gaussian Processes - 2. Use Line Features
=, prepare for Gaussian Process Regression In this bachelor thesis Implement
3D face registration using Gaussian Processes and Line Features. One part of
the problem is to sample equidistant 3D points from 2D line features marked
on images of a 3D face scan. These line features should then be used as
an additional input to a registration algorithm which is based on Gaussian
Process Regression. The aim is to build a pipeline which starts off with the
raw scan data as well as the landmarks and line features. The feature points
are used to register the mean face of the MM/BFM (Basel Face Model) on
to/with the raw scan thereby obtaining a fully defined and textured 3D model
representation of the face in 3D. Registration is the technique of aligning to
objects using a transformation, in this case the registration is performed by
adding displacements to every points in the mean face model. A model is
represented as vector $N \times d$. What is a model? A vector representation of a 3D
scan? For the morphing a Posterior Shape Model is used in combination with
a Gaussian Process. Image registration is a process of aligning two images
into a common coordinate system thus aligning.
(gaussian process + line features for accurate, reproducible registration)

1.2 Review Literature

2. Definition of terms (morphable model, 3D face registration, Gaussian Process regression, posterior shape models) 3. Review of literature (papers)

Chapter 2

3D Model Building

This chapter contains a short overview on how to build a generative 3D face model from 3D face scans. (because this is what should eventually be achieved with this thesis). The term generative describes the fact that with a linear face model, arbitrary, feasible faces can be generated by combining and scaling faces that are part of the model.

2.1 3D Morphable Model

The 3D Morphable Model (3DMM) published by Blanz and Vetter in 1999 (bib) is such a generative 3D face model. It is a linear model built on a set of faces parametrized by coefficients $\vec{\alpha}$ and a set of textures parametrized by coefficients $\vec{\beta}$

$$\mathcal{S}(\alpha) = s + S\alpha \quad \mathcal{T}(\beta) = t + T\beta \quad (2.1)$$

The average of the faces $s \in \mathbb{R}^{3N}$ and of the textures $t \in \mathbb{R}^{3N}$ are used to fit two multivariate normal distributions to the whole dataset of faces. The covariance matrices are defined over the differences between each face and the mean face, the same applies to the textures. We get two distributions. A set of faces parametrized by coefficients α , set of Textures parametrized by coefficients β

(2.2)

Fit multivariate normal distribution to data set, based on average of faces and textures. Build covariance matrices over differences between the mean and face samples in surface and texture. => two distributions. Perform PCA to get orthogonal basis system In the MM three subspaces are morphed independently

Eine Gruppe (G,) heisst abelsch [abelian] oder kommutativ wenn $ab = ba$ gilt fr alle $a, b \in G$.

2.3 Building a Model from the registered data (short)

114 2.2 Prerequisite Data

115 image with landmarks and line features a short overview what data we have
116 given

117 Facial Scans: face scans given as point clouds The data we have given
118 is a set of about 300 face scans that have had a set of key points marked.
119 Furthermore important and detailed regions like the eyes, ears and lips have
120 been marked by contour lines known as line features. The scans have been
121 obtained with a scanner. The surface is very detailed, however the eyes and
122 the nostrils are not recorded. From these scans we want to create fully textured
123 3D faces, which can be used to build a new face model.

124 Mean Face: The mean face has been derived from a collection of 100 male
125 and 100 female 3D face models.

126 2.3 Finding Correspondences

127 WE WANT POINT TO POINT CORRESPONDENCE BETWEEN THE
128 TWO FACES in general: point to point correspondence between to images
129 Are scans already in semantical correspondence? No semantical correspon-
130 dence FINDING CORRESPONDENCE IS EXACTLY THE AIM OF REG-
131 ISTRATION =_ HAVING SAME POINTS AS CLOSE TO ONE ANOTHER
132 AS POSSIBLE Now in order to obtain a 3D representations of the face we
133 need to transform the mean face so that it fits a particular 3D face scan. To
134 find the transformation, however, we first have to find feature points in both
135 3D representations which correspond to the same semantical structure. Pre-
136 vious work has shown that point landmarks are not sufficient to preserve the
137 level of detail which is imminent in the regions of the eyes, ears and lips and
138 that the computed transformations are not able to preserve these regions. For
139 this reason, additional line features have been introduced. In order to relate
140 these

141 How registration works so far

142 What we want to change

Chapter 3

Gaussian Processes in 3D Face Registration

As described briefly in the introduction, the first of our two objectives is to build a face registration pipeline. In achieving this we use an algorithm which can handle arbitrary shapes for the registration of the 3D faces. It is derived from a stochastic process, more specifically a vector-valued Gaussian process or Gaussian random field. In this chapter we deal with the theory necessary for understanding the functionality of the registration pipeline. To begin with, we recapitulate the definition of stochastic processes and extend it to the definition of Gaussian processes. In the next step, we then delve into Gaussian process regression and finish by applying a vector-valued Gaussian process to our problem of 3D face mesh registration.

3.1 Stochastic Processes

In probability theory a stochastic process consists of a collection of random variables $\{X(t)\}_{t \in \Omega}$ where Ω is an index set. It is used to model the change of a random value over time. The underlying parameter time is either real or integer valued. In our case, however, we use a collection of vector-valued random variables with indices in the \mathbb{R}^3 space because we want to model deformations of the vectors on the faces' surfaces. This generalization of a stochastic process - which can handle multidimensional vectors - is called a random field. Defining the random variables on an index set in an n-dimensional space, allows for spatial correlation of the resulting values, which is an important aspect of the algorithm discussed later on.

3.2 Gaussian Processes

A Gaussian process is a stochastic process in which each finite collection $\Omega_0 \subset \Omega$ of random variables has a joint normal distribution. More formally, we define the collection of random variables $\{X(t)\}_{t \in \Omega}$ to have a d-dimensional normal distribution if the collection $\{X(t)\}_{t \in \Omega_0}$ - for any finite subset Ω_0 - has a joint $d \times |\Omega_0|$ -dimensional normal distribution with mean $\mu(\Omega_0)$ and covariance $\Sigma(\Omega_0)$. If $\Omega \not\subseteq \mathbb{R}$ holds, the process is a Gaussian random field, which holds true for our case, because we use an index set $\Omega \subseteq \mathbb{R}^3$. In the

175 further proceedings the term “Vector-valued Gaussian Processes” will be used
 176 to refer to Gaussian random field” will be used to refer to Gaussian random
 177 fields.

178 An alternative way of viewing a Gaussian process is to consider it as a
 179 distribution over functions. This allows us to look for inference in the space of
 180 these functions given a dataset, specifically to find the deformation function
 181 given a 3D face mesh. Each random variable now yields the value of a function
 182 $f(x)$ at a location $x \in \mathcal{X}$ in the index set of possible inputs. We now denote
 183 the index set by \mathcal{X} to stress that we are ceasing to discuss Gaussian processes
 184 defined over time. In this function-space view a Gaussian Process at location
 185 x is thus $f(x) \sim GP(\mu(x), k(x, x'))$ defined by its mean $\mu : \mathcal{X} \rightarrow \mathbb{R}$ and
 186 covariance $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ functions which in turn are defined over the set
 187 of input vectors. With $\mu(\mathcal{X}) = (\mu(x))_{x \in \mathcal{X}}$ and $\Sigma(\mathcal{X}) = (k(x, x'))_{x, x' \in \mathcal{X}}$ we
 188 obtain the full distribution of the process $GP(\mu(\mathcal{X}), \Sigma(\mathcal{X}))$. For the purpose
 189 of simplifying calculations we may assume that every random variable has zero
 190 mean without a loss of generality. When modeling a deformation field with a
 191 Gaussian process this circumstance implies that the expected deformation is
 192 itself zero.

193 **Covariance Functions** The key feature of a Gaussian Process is its covari-
 194 ance function also known as “kernel”. It specifies the covariance $\mathbb{E}[f(x)f(x')]$
 195 between pairs of random variables for two input vectors x and x' , allowing
 196 us to make assumptions about the input space by defining the spatial co-
 197 dependency of the modelled random variables. Note that when assuming zero
 198 mean we can completely define the process’ behaviour with the covariance
 199 function.

200 A simple example of a covariance function is the squared exponential covari-
 201 ance function, defined by $cov(f(x), f(x')) = k(x, x') = \exp(-\frac{(x-x')^2}{2l^2})$. (deriva-
 202 tion Rasmussen et al. p.83) *still to be continued and refined...*

203 It is possible to obtain different prior models by using different covariance
 204 functions. In our case, we use a stationary (x-x, invariant to translation),
 205 isotropic exponential covariance function - Squared Exponential Covariance
 206 Function (p. 38)

207 **Gaussian Process Prior** The specification of the covariance function im-
 208 plies that a GP is a distribution over functions. To illustrate this one can
 209 draw samples from a prior distribution of functions evaluated at any number
 210 of points, X_* . The Gaussian Process Prior is solely defined by the covariance
 211 matrix made up of the covariances of the input points.

$$\Sigma(X_*) = \begin{bmatrix} k(x_{*1}, x_{*2}) & \cdots & cov(f(x_{*1}), f(x_{*n})) \\ \vdots & \ddots & \vdots \\ cov(f(x_{*n}), f(x_{*1})) & \cdots & cov(f(x_{*n}), f(x_{*n})) \end{bmatrix} \in \mathcal{M}^{|X_*| \times |X_*|} \quad (3.1)$$

212 A sample is a random Gaussian vector $f_* \sim \mathcal{N}(0, \Sigma(X_*))$ containing a
 213 function value for every given input point. Plotting random samples above
 214 their input points is a nice way of illustrating that a GP is indeed a distribution
 215 over functions, see figure 3.1. The GP Prior forms the basis for inference in
 216 Gaussian Process Regression.

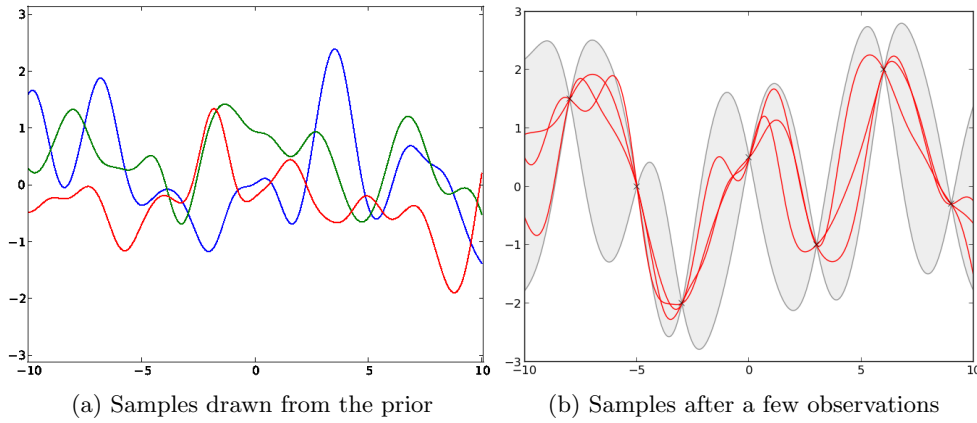


Figure 3.1: In figure a) three functions have been drawn from the GP prior, each has a sample size of 1000 points. Figure b) again shows three functions, but this time the prediction incorporates the information of random values observed at seven points in the input space

217 **Vector-valued Gaussian Processes** In order to use Gaussian processes
 218 to model deformation fields of three dimensional vectors as intended, there is
 219 the need for a generalization of the above definition from the function-space
 220 view. The random variables $X_1, X_2, \dots, X_k, \dots, X_n$ are now d-dimensional
 221 vectors, yielding a covariance function of the form $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}^{d \times d}$ and
 222 $k(x, x') = \mathbb{E}[X_k(x)^T X_k(x')]$. *Should this paragraph be continued?*

223 3.3 Gaussian Process Regression

224 The task of registering two 3D face meshes can be treated as a regression
 225 problem in which the goal is to predict the deformation of all floating mesh
 226 points, given the displacement of the landmarks present in both meshes. Trying
 227 to fit an expected function - be it linear, quadratic, cubic or nonpolynomial
 228 - to the data is not a sufficiently elaborated approach to our problem. Using
 229 a Gaussian Process disposes of the need to describe the data by a specific
 230 function type, because the response for every input point is now represented
 231 by a normally distributed random value, in turn governed by the specification
 232 of the covariance function.

233 Key assumption: data can be represented as a sample from a multivariate
 234 gaussian distribution P

235 **Regression Problem** Assume a training set $\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$
 236 where $x \in \mathbb{R}^d$ and y is a scalar output or target. Later on, in the case of the
 237 training set consisting of landmarks, a Vector-valued Gaussian Process must
 238 be used, because y is then also a vector $y \in \mathbb{R}^d$. The task is now to infer
 239 the conditional distribution of the targets for yet unseen inputs and given the
 240 training data $p(\mathbf{f}_* | \mathbf{x}_*, \mathcal{D})$

241 **Noise-free Prediction** First we assume the observations from the training
 242 data to be noise-free so that we can fix the training data to these observations

243 \mathbf{y} without complicating the model. The joint prior distribution with training
 244 \mathbf{f} and test \mathbf{f}_* outputs indicated is the following:

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}_* \end{bmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} \Sigma(X) & \Sigma(X, X_*) \\ \Sigma(X_*, X) & \Sigma(X_*) \end{bmatrix} \right) \quad (3.2)$$

245 We can now obtain the posterior samples illustrated in 3.1 b) by condition-
 246 ing the above joint Gaussian prior distribution on the observations $\mathbf{f}_* | \mathbf{f} = \mathbf{y}$
 247 which results in the following distribution:

$$\mathbf{f}_* | X_*, (X, \mathbf{f}) \sim \mathcal{N} \left(\Sigma(X_*, X) \Sigma(X)^{-1} \mathbf{f}, \Sigma(X_*) - \Sigma(X_*, X) \Sigma(X)^{-1} \Sigma(X, X_*) \right) \quad (3.3)$$

248 Later on, we will extend this definition to 3-dimensional inputs and out-
 249 puts.

250 **Prediction with Gaussian Noise Model** In most real world applications
 251 as is the case for the problem we will look into later, however, observations
 252 from the training data are not free of noise. The landmarks clicked on the
 253 3D face meshes, for example, can never be marked at the exact same feature
 254 location. These circumstances call for the incorporation of a noise model.
 255 We specify a simple additive i.i.d Gaussian noise model $y = f(x) + \varepsilon$ where
 256 $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ for every input vector \mathbf{x} . In section ?? the variances will be
 257 varied for every sole landmark. For now it is enough to add the variance of
 258 the noise model to the covariance of the training.

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{y}_* \end{bmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} \Sigma(X) + \sigma^2 \mathcal{I}_{|X|} & \Sigma(X, X_*) \\ \Sigma(X_*, X) & \Sigma(X_*) \end{bmatrix} \right) \quad (3.4)$$

The distribution - now conditioned on the noisy observations - is thus

$$\mathbf{y}_* | \mathbf{f} = \mathbf{y} \sim \mathcal{N}(\bar{\mathbf{y}}_*, \Sigma(\mathbf{y}_*)) \quad (3.5a)$$

where the mean depends on the observed training targets

$$\bar{\mathbf{y}}_* = \Sigma(X_*, X) (\Sigma(X) + \sigma^2 \mathcal{I}_{|X|})^{-1} \mathbf{y} \quad (3.5b)$$

whilst the covariance depends only on the input points

$$\Sigma_* = \Sigma(X_*) - \Sigma(X_*, X) (\Sigma(X) + \sigma^2 \mathcal{I}_{|X|})^{-1} \Sigma(X, X_*) \quad (3.5c)$$

259 *Conclusion, how does this help us to proceed?*

260 3.4 Application to 3D Face Meshs

261 *deformation field bold instead of calligraphic?* We strive to predict a
 262 deformation field $\mathcal{D} : \mathcal{M} \subset \mathbb{R}^3 \rightarrow \mathbb{R}^3$ which assigns a displacement vector to
 263 every vertex in a triangulated mesh. The mesh in question is called a floating
 264 or moving mesh \mathcal{M} . Adding the displacement field to the moving mesh should
 265 then render a mesh corresponding as closely as possible to a target mesh \mathcal{T}
 266 and thereby performing a registration / should provide an accurate mapping
 267 on to the target mesh. The mean mesh of the Morphable Model serves as the
 268 moving mesh which we want to register with multiple triangulated meshes of
 269 scanned target faces.

Reference Mesh Prior As defined by the deformation field the output the regression problem is in \mathbb{R}^3 calling for the use of a Vector-valued Gaussian Process with random variables $u \subseteq \mathbb{R}^3$. After the reference and target have been aligned ?? a Vector-valued Gaussian Process can be initialized by defining the prior over all vertices of mean mesh. For this purpose the covariance function has to be redefined to handle 3-dimensional vectors. *Prior consists of smooth deformations of the mean face*

$$k \left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}, \begin{bmatrix} x'_1 \\ x'_2 \\ x'_3 \end{bmatrix} \right) = xy^T \in M^{3 \times 3} \quad (3.6)$$

Each covariance entails 9 relationships between the different components of the vectors, yielding a 3×3 matrix. The covariance matrix then grows to become:

$$\Sigma_{\mathcal{X}} = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{bmatrix} \in M^{3n \times 3n} \quad (3.7)$$

The mean mesh is defined by a set of vectors $\mathcal{X} \in \mathbb{R}^3$ and a set of landmarks $L_{\mathcal{M}} = \{l_1, \dots, l_n\} \subset \mathbb{R}^3$. The mean vector μ is made up of the component-wise listing of vectors so that is has dimensionality $3n$. Setting the whole mean vector to zero, as discussed before, implies a mean deformation of zero and makes perfect sense in this setting, because we are modelling deformations of the mean face surface. The prior distribution over the mean face mesh is therefore defined as

$$\mathcal{D} \sim \mathcal{N}(\mathbf{0}, \Sigma_{\mathcal{X}}) \quad (3.8)$$

270 meaning that a deformation field can be directly drawn as a sample from the
 271 prior distribution of the vertices of the mean mesh. *show two or three*
 272 *samples of prior here, next to mean mesh*

Reference Mesh Posterior The target landmarks also consist of a set $L_{\mathcal{T}} = \{l_{\mathcal{M}1}, \dots, l_{\mathcal{M}N}\} \subset \mathbb{R}^3$. Fixing the prior output to the deformation vectors $\mathcal{Y} = \{\mathbf{t} - \mathbf{m} | \mathbf{t} \in L_{\mathcal{T}}, \mathbf{m} \in L_{\mathcal{M}}\}$ defined by the distance between the reference and target landmarks and assuming additive i.i.d Gaussian noise the resulting posterior distribution is

$$\begin{bmatrix} \mathcal{Y}_{\varepsilon} \\ \mathcal{D} \end{bmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{bmatrix} \Sigma(\mathcal{Y}) + \sigma^2 \mathcal{I}_{3|\mathcal{Y}|} & \Sigma(\mathcal{Y}, X_*) \\ \Sigma(X_*, \mathcal{Y}) & \Sigma(X_*) \end{bmatrix} \right) \quad (3.9)$$

273 *Is this a correct definition for the distribution?*

274 The deformation model is now rendered fixed at certain landmark points
 275 in the target mesh and the goal is to find valid deformations through the set of
 276 fixed targets, analogous to the case of eq. 3.5a. The posterior model is defined
 277 as the joint distribution of all mean mesh points and the mean landmarks,
 278 conditioned on the output deformation vectors for every mean landmark with
 279 added noise.

$$\mathcal{D} | \mathcal{X} \rightarrow \mathcal{Y}_{\varepsilon}. \quad (3.10)$$

280 e now have defined a distribution over our mean face mesh. The variance of the
 281 gaussian kernel can thereby be described as a smoothing parameter P *mean*
 282 *is now max a posteriori solution*

283 Sampling the conditional distribution creates deformed 3D surfaces of the
 284 mean mesh which are fixed at the target landmarks. *show images of mean,*
 285 *prior and posterior with added landmarks*

286 3.5 Fitting & Optimization

Due to the fact that the posterior model is fixed at the target landmarks it is possible to perform the registration by drawing a shape from the posterior model. The sample, however, has to be optimized according to the shape of the target face. In order to find the deformation corresponding to the optimal fit d_* a linear optimization with the posterior process as a constraint is employed/ regularization term. (small lambda) a bit of the posterior mean)

$$d_* = \arg \min_{d \in \mathcal{D}} L[O_T, O_M \circ d] + \lambda R[d] \quad (3.11)$$

Minimizing a loss function L - mean square distance for example - on the target and the deformed mean provides a feasible deformation field. \mathcal{D} denotes the space of possible deformations in the posterior model. The information needed is held by the covariance matrix of the posterior process. However, using Mercer's theorem (in short what it does, yada yada yada) a basis of functions corresponding to all possible deformations can be extracted. The kernel is thus described as a linear model of these basis functions. *whole matrix or simple kernel? refer to paper "a unified formulation of statistical model fitting and non-rigid registration"*

$$k(x, x') = \sum_{i=1}^{\infty} \lambda_i \phi_i(x) \phi_i(x') \quad (3.12)$$

λ_i are the eigenvalues and ϕ_i the eigenvectors of K . They denote the deformation directions while the eigenvalues \dots . We are looking for a finite linear combination of eigenvectors that form a deformation field with $\exists \alpha_1 \dots \alpha_n \in \mathbb{R}$ as linear parameters.

$$f(x) = \sum_{i=1}^n \alpha_i \lambda_i \phi_i(x) \quad (3.13)$$

287 f $\text{GP}(0, K)$ we take our gaussian process \hat{f} — $x=y$, *ask Marcel for a*
 288 *helping hand with the theory?*

Next, we want to minimize residuals for the whole of our surface according to optimal parameter values α_i

$$\arg \min_{\alpha \in \mathbb{R}^n} \sum_{x_i \in \mathcal{T}_R} (f(x_i) - \phi_T(x_i))^2 \quad (3.14)$$

where $f(x_i)$ is the deformation function and $\phi_T(x_i)$ returns the nearest point on the target mesh. Yields the overall loss function Φ_L

$$\Phi_L(f(x_i) - \phi_T(x_i)) \quad (3.15)$$

289 The eigen vectors - which are deformation vectors defining a deformation
 290 for every model vertex - of the covariance matrix define a basis space? Shape
 291 Modell = select best eigenvectors via PCA in order to simplify computation.
 292 = vorstellen wie wenn mehrere Wellbleche durch die Target "landmarks
 293 gelegt werden und dann mit bestimmten parametern alpha zwischen ihnen
 294 interpoliert wird Alternative way to understand basis functions for gaussian
 295 process: sample from the GP(0, K) and then build a linear model from the
 296 functions, $f(x) = \sum_i \alpha_i \phi_i(x)$ Posterior Distribution of Landmarks
 297 Defining the Gaussian Process Posterior Distribution - Landmarks (Referenz
 298 deformieren From Gaussian Processes to Shape Models = by selected principal
 299 components of the covariance matrix

300 3.6 Robust Loss Functions

robust against outliers the Alignment of the mean face mesh and the target mesh causes overlaps on the forehead, the side of the head and the neck. Using a simple Mean Square error between the reference and target mesh for optimization penalizes the overlapping regions with a strong gradient and therefore causes strong distortions. Our approach to tackling this problem was to try out a range of different robust estimators, namely the Tukey, Huber, and Fair estimators. (table with formulas?) The advantage is that these estimators are less sensitive to outliers, reducing the artefacts of registration considerably. However, as can be seen from the formulas, these techniques require finding appropriate parameters first which produce reasonable/acceptable visual results Fair

$$\rho(x) = c^2 \left[\frac{|x|}{c} - \log\left(1 + \frac{|x|}{c}\right) \right] \quad (3.16a)$$

$$\psi(x) = \frac{x}{1 + \frac{|x|}{c}} \quad (3.16b)$$

Huber

$$\rho(x) = \begin{cases} \frac{x^2}{2} & \text{if } |x| < k \\ k(|x| - \frac{k}{2}) & \text{if } |x| \geq k \end{cases} \quad (3.17a)$$

$$\psi(x) = \begin{cases} x & \text{if } |x| < k \\ k \operatorname{sgn}(x) & \text{if } |x| \geq k \end{cases} \quad (3.17b)$$

Tukey

$$\rho(x) = \begin{cases} \frac{c^2}{6} \left(1 - \left[1 - \left(\frac{x}{c} \right)^2 \right]^3 \right) & \text{if } |x| \leq c \\ \frac{c^2}{6} & \text{if } |x| > c \end{cases} \quad (3.18a)$$

$$\psi(x) = \begin{cases} x \left[1 - \left(\frac{x}{c} \right)^2 \right]^2 & \text{if } |x| \leq c \\ 0 & \text{if } |x| > c \end{cases} \quad (3.18b)$$

301 **Rigid Alignment** This part is the theoretical part, alignment can follow
 302 in the Pipeline Therefore, we first have to perform a rigid transformation to
 303 align the meshes according to the feature/ landmark points.

Chapter 4

Registration Pipeline using Line Features

In this chapter we follow up on the definition of the Vector-valued Gaussian Model for 3D Face Registration by describing the registration pipeline built to put this concept into practice. The pipeline is of a sequential nature, where in each step the output of a data processing unit is the input for the next step. To enhance the registration outcome of this pipeline we use contour lines of key regions of the face.

4.1 Line Features

Definition of Line Features

For every scan we want to register, 8 contours have been marked on three images of the face - taken from the front, the left and the right of the face - with a special GUI for marking points and lines on images. These contours depict the eyebrows, eyes, ears and lips of a face and we call them “line features”. They are made up of a set of segments, each of which is modelled with a **Bézier curve** (parametric curve frequently used in computer graphics, bernstein basis polynomials, used for modelling smooth curves) of varying order. Due to the nature of the objects depicted, there are open as well as closed curves.

$$B(t) = \sum_{i=0}^n (1-t)^{n-i} t^i P_i \quad (4.1)$$

The line features are saved in explicit files along with the face mesh of the scan.

Why use Line Features for Registration?

Line features serve the purpose of augmenting the quality of registration by initiating it with a larger set of corresponding points (points which are on the lines). They are used to mark complex regions of the eyes, i.e. the eyes, ears etc., so that the registration process produces an accurate mapping of the contours of these organs which would otherwise not be possible. Areas containing “curves” have a dense abundance of points/parameter changes, while straight areas only have scarce points.



Figure 4.1: Line Features of the left eyebrow and eye, consisting of bézier curves defined by visible control points (white)

4.2 Sampling 3D Points from 2D Line Features

In order to be able to use line features in the Vector-valued Gaussian Model, they have to be sampled at discrete intervals resulting in a set of additional landmarks $L_{Add} = \{l_1, \dots, l_N\}$. These define the mapping $\Omega : L_{Add\mathcal{M}} \rightarrow L_{Add\mathcal{T}}$ of the contours - describing the different important features present in the faces - in the mean face mesh on those of the target face mesh. In order for the mapping Ω to be plausible, it is essential for the curves to have equidistant parametrization so that when curves undergo sampling of N points, these N points are all at equal parametric intervals.





Figure 4.2: Mapping of equidistant samples of **ear** line features from the reference (left) on to the target (right)

334 Arc Length Parametrization

335 The first problem which becomes apparent when trying to sample the line fea-
 336 tures is that the bézier curve segments don't allow for equidistant parametriza-
 337 tion, because the underlying parameter $t \in \mathbb{R}$ is not linear in respect to the
 338 length of the curve. The growth of the parameter of a bézier curve is instead
 339 dictated by velocity.



340 *add another point to the right ear, so there are 11.*

341 Consequently, the imperative must be to evaluate the curves based on their
 342 arc-length, which is defined as the length of the rectified curve, instead. The
 343 underlying parameter must then correspond - at every point of the curve -
 344 to the ratio of the curve length that has been traversed and the total curve
 345 length.

346 **In theory** It is possible to get the arc length $L(t) = \int_{t_0}^{t_1} |C'(t)| dt$ for given
 347 parameters t_0, t_1 where $C'(t)$ is the derivative of the curve $C : t \in [0, 1] \rightarrow \mathbb{R}^2$.
 348 What we are in need of, however, is a reverse mapping from the length of a
 349 fraction of the curve to the curve parameter $t = L^{-1}(l)$. This mapping can
 350 of course be derived analytically, but it is far easier to implement it using a
 351 numeric approximation.

352 **In practice** As we are not in need of a subpixel accurate resolution, we can
 353 skip the formal math and use a lookup table to compute the arc-length. First,
 354 we calculate $n=1000$ points on each segment the curve - made up of bézier
 355 curves using the normal parameter t . For each point we save the euclidean
 356 distance from the origin of the segment into a new slot in the lookup array.
 357 We get the euclidean distance for one point by summing up the distance to
 358 the predecessor/preceeding points and its distance from the origin.

draw a line with a few segments draw curve with points just off and

359
 360 *lookup table beneath* In effect, we are provided with have a lookup array
 361 that contains the approximated distances of a large number of points from the
 362 origin of a curve segment. Assembling the segments' lookup arrays gives us
 363 the overall array for the curve with the last value presenting the arc-length of

the whole curve.
 Second, finding points on the curve according to a linear parameter governed by the amount of points that we want to parametrize the curve with is quite easy. The curve can easily be sampled by computing the length of parametric intervals $\frac{L}{N}$ for a specified number of points N to be sampled. $l = k \cdot \frac{L}{N}$ returns the current length of the curve for the sampling point of index k , where $k = [0, \dots, N]$ for open curves and $k = [0, \dots, N-1]$ for closed curves. Then we simply perform a binary search on the lookup table (to get largest value smaller than n ?) for this distance. We choose the index that returns the exact length we specified or the index with the next smaller length. The coordinates of the point with this index t are now the coordinates we use for the sampling point. We compute the distance we want to travel the curve using the length of equidistant sections and the point we want to get. *reference lines in text above?*

Listing 4.1: Equidistant Sampling

```

378 void getEquidistantPoints(int numSampleSegments = 20) {
379     // static members:
380     // arcLookup - lookup table
381     // totalLength - total arc-length of curve
382     // auxiliaryPoints - ??? exact definition
383
384     if(arcLookup.size() == 0) return;
385     int pointsToDraw = numSampleSegments+1;
386     if(closed) pointsToDraw--;
387
388     T sectionLength = totalLength/numSampleSegments;
389
390     for(size_t i=0; i < pointsToDraw; ++i) {
391         T progress = i*sectionLength;
392         // perform c++ binary search on lookup table
393         int low = 0;
394         int currIndex = 0;
395         int high = arcLookup.size()-1;
396         T currPieceLength;
397
398         while(low < high) {
399             currIndex = low + (high - low)/2;
400             currPieceLength = arcLookup[currIndex];
401             if(currPieceLength < progress) {
402                 low = currIndex+1;
403             } else {
404                 high = currIndex;
405             }
406         }
407         // currPieceLength is now >= progress
408         if(currPieceLength > progress) {
409             currIndex--; // currPieceLength is now < progress
410         }
411         equidistantPoints.push_back(auxiliaryPoints[currIndex]);
412     }
413 }

```

Mesh Projection of Sampled Points

Having implemented arc length parametrization it is possible to draw an arbitrary amount of samples $x \in \mathbb{R}^2$ from the line features. They are thereby defined as a set of points $S \subset \mathbb{R}^2$. Our goal is, however, to have these additional landmarks describing the features on the mesh itself and not a 2-dimensional snapshot. We therefore need to use the camera calibration and some computer graphics to project the sampled points onto a face mesh for each line feature we want to obtain.

Projection from camera to mesh? How? Knothe

In the previously used registration method, a large number of points was used for each curve. These points were, however, not projected directly on to the 3/4 shells of the mesh. Instead their location was constrained by computing a 1-dimensional band of points before and after their approximate position, seen from the origin.

Get direction of 3D representation of a curve, compute distance from origin to mesh, normalize to direction.

Compute distances from origin for mesh vertices dot product (direction of point, for every vertex: direction of vertex) dot product: 1 for similar directions, 0 for perpendicular directions. Angle value \angle .9999: mindist, maxdist are updated with the distance value of the distance vector to the nearest vertex on to the direction of the actual 3D representation of the point

In the previous registration method implemented by Dr. Brian Amberg, the points from line features constrained to 1D and are then projected on to the 3 shell meshes with the program points_from_surface. That is how the feature points are generated. shells from the scanner are cleaned points are marked on the 3 images to the front, left and right of the person explain what program does, latex sketches: camera calibration is done by the scanner? meshes and camera settings are loaded into the software. For a lot of points per curve the direction their 3d representation is computed and their distance, normalized to give direction, from the origin is saved. Distances and directions are further computed for all vertices in the mesh. Now for every point the dot product of its direction is formed with direction of every vertex in the mesh. Remember, the dot product results in 1 for similar directions and 0 for perpendicular directions. For an angle value larger than .9999 the parameters min_dist or max_dist are updated with the distance value of the projection of the distance vector to the nearest vertex on to the direction of the actual 3D representation of the point. min(min_dist, d), max(max_dist, d), so that at the end min_dist contains the distance to the point on the line nearest to the camera and max_dist the distance to the point farthest away from the camera. now min_dist is multiplied with a value \angle 1 and max_dist with a value \angle 1. then a band of points is computed perpendicular to the line along the direction at each computed point that cuts through the mesh and serves as a horizontal constraint for the points to lie after the registration. picture of dot product

Listing 4.2: Point Projection

```

457   for (size_t l=0; l<templates.size(); ++l) {
458       if (!templates[l].isSet)
459           continue;
460       {
461           vector<f3Vector> meanEqPoints3d;
462           templates[l].evaluate(512);
463           // Find min_dist and max_dist for this template
464           /* COMPUTE EQUIDISTANT LINE POINTS AND WRITE THEM TO FILE */
465
466           // computing 1000 points per curve segment
467           templates[l].curve.initializePoints();
468           // create arc length lookup table computed over all initialized
469           // segment points
470           templates[l].curve.approxTotLength();
471           // equidistant sampling
472           templates[l].curve.getEquidistantPoints(numPoints);
473           vector<d2Vector> eqPoints = templates[l].curve.equidistantPoints;
474           vector<d3Vector> eqDirs;
475
476           for (size_t i=0; i<eqPoints.size(); ++i) {
477               // compute direction vector of 3d representation of points on
478               // curve from the point of origin
479               auto point = eqPoints[i];

```

```

480         d3Vector dir= -O + C.imageToWorld(point);
481         dir /= dir.normL2();
482         eqDirs.push_back(dir);
483     }
484     vector<double> selectedVecDist;
485     // go over all directions of points on the line template and
486     // compute the dotproduct with the current mesh vertex
487     // direction
488     for (size_t p=0; p < eqDirs.size(); ++p) {
489         const d3Vector &dir = eqDirs[p];
490         // save distances along the directions of near vertices and
491         // angles for every point
492         vector<double> remDistances;
493         vector<double> remAngles;
494         for (size_t i=0; i < meshes.size(); ++i) {
495             for (size_t j=0; j < meshes[i].vertex.size(); ++j) {
496                 // compute direction from origin for every vertex in
497                 // the mesh
498                 d3Vector vert_dir = (d3Vector(meshes[i].vertex[j]) -
499                     O);
500                 d3Vector vert_dir_n = vert_dir / vert_dir.normL2();
501
502                 double a = vert_dir_n.dot(dir);
503                 // if direction likeness is bigger than 99.99%
504                 if (a > 0.9999) {
505                     // projection of distance vector of mesh vertex
506                     // onto direction of 3d representation of true
507                     // point on curve segment
508                     // project vert_dir onto dir and
509                     double dist = vert_dir.dot(dir);
510                     remDistances.push_back(dist);
511                     remAngles.push_back(a);
512                 }
513             }
514         }
515         // choose distance via best angle match, PROBLEM: holes in
516         // mesh
517         if (remAngles.size() > 0) {
518             int index = std::max_element(remAngles.begin(), remAngles
519                 .end()) - remAngles.begin();
520             selectedVecDist.push_back(remDistances[index]);
521         } else {
522             selectedVecDist.push_back(0.0);
523         }
524     }
525     // save directions to equidistant points on all line features in
526     // 3D
527     for (size_t i=0; i < eqPoints.size(); ++i) {
528         d3Vector dir = -O + C.imageToWorld(eqPoints[i]);
529         dir /= dir.normL2();
530         float a = 0.5f;
531         f3Vector point;
532         if (selectedVecDist[i] != 0) {
533             f3Vector tmp(O + selectedVecDist[i] * dir);
534             point = tmp;
535         }
536         meanEqPoints3d.push_back(point);
537     }
538 }

```

539 **Modification** *up close image of eye holes of face scan* The modifica-
540 tion we introduced, was solely to select the mesh vertex with the highest
541 similarity of direction to the 3D representation of a 2D line feature sample.
542 Due to the areas of the mesh around the ears and the eyes containing large
543 holes the projected sample points from the line features can be off target,
544 especially if a large number of points is sampled for every curve. This cir-
545 cumstance leads to the sampled line features being represented by more of a
546 point cloud, for example around the eyes, (which is not distinguishable as a
547 line) instead of clearly denoting a contour line. The direction of the vertex is
548 used to find a point –; this distorts the shape (position of sampled points) of
549 the line. On different data sets the performance of the projection of the line
550 features for a large number of samples varied enormously. *Compute some*
551 *landmarks with 30 samples* Using only 5-10 sample points per curve some

552 datasets rendered near perfect results on a “control sample”. *Compile list*
 553 *of datasets* However, as long as the method is dependent on the data from
 554 the scans - the size of the holes in the meshes - it lacks generality and generality
 555 is exactly the basis for feasible and reproducible registration results.

```

556      /**
557      * Returns the position in world coordinates lying on the focal plane
558      * which is corresponding to a pixel coordinate. The camera ray is
559      * c.imageToWorld(v_i) - c.origin()
560      */
561      inline
562      t3Vector<T> imageToWorld(const t2Vector<T> &v_i) const {
563          /// Pixel to Distorted Image Plane
564          //
565          /// Offset
566          const t2Vector<T> v_i_o = v_i - C;
567          /// Scale
568          const t2Vector<T> v_d(v_i_o.x/sx*d.x, v_i_o.y*d.y);
569
570          /// Distorted Image to Pinhole
571          const t2Vector<T> v_p = v_d * (1.0 + k1*v_d.normL2sqr());
572
573          /// Pinhole to Camera
574          const t3Vector<T> v_c(v_p.x, v_p.y, f);
575
576          /// Camera to World
577          const t3Vector<T> v_w = Rinv * (v_c - t);
578          return v_w;
579      }

```

580 4.3 Preparing the Mean Mesh

581 Rendering, marking = Projection On top of that, another problem occurred,
 582 because the mean face mesh of course doesn't have any line features projected
 583 on to it either. Rendering, marking line features, projecting back possible,
 584 because we know direction

585 However, it contains about 60 feature points manually clicked, which are
 586 not present in newly scanned datasets. Eliminate the ones, which are not
 587 clicked on scans

588 **Output** pipeline specifications

589 4.4 Rigid Mesh Alignment

590 simple rigid transformation of the scanned face onto the mean, transformation
 591 computed from landmark vectors. To begin the registration we first have to
 592 align the two meshes. The floating mesh has to be clipped at the neck and
 593 around the ears where the scanner has left artifacts. Furthermore the mouth
 594 cavity of the mean face has to be removed. We then selected the 11 feature
 595 points present in the floating mesh in the mean face from the abundant 60.
 596 To achieve this we wrote a python script loading the feature point files. A
 597 feature point is described by its 3D coordinates, a visibility parameter in the
 598 range [-1,1] and a label denoting its exact location (mouth.inner.upper). All
 599 we had to do now was to create dictionaries label : (x,y,z) and to compare
 600 them for labels. Then we passed the resulting point correspondencies to the
 601 python vtk api for the mean of computing a transformation comprised of
 602 simple translation and rotation (no scaling, only 3 point correspondencies
 603 needed). Note, we are not trying to map the meshes on to one another here.
 604 We are simply trying to align them through the use of the feature points. The

605 computed transformation we applied to all points in the floating mesh. The
 606 resulting mesh was written to a file and then opened in paraview. We now
 607 had the meshes in a position from where we could start the actual mapping.
 608 The mean face was broader in shape than the scan and was perfectly coated in
 609 texture for the simple reason that hours of manual labour have been invested
 610 to render this important piece of data a perfect reference. Now in order to
 611 receive a perfect mapping of the floating mesh on to the mean/reference mesh
 612 we have to allow for 3 degrees of freedom, that is in all 3 dimensions x,y
 613 and z, for every pixel in the floating mesh except for the reference points we
 614 have used as correspondencies. The parameters having the most influence to
 615 the mapping will be those specified in the constraints we introduced into the
 616 equation via regularization. The idea behind the use of sampled points from
 617 the line features was to have more point correspondencies in complex regions
 618 as for example the eyes and the ears where there is a great abundance of
 619 pixels and the algorithm isnt likely to create a flow field which is accurate not
 620 enough to describe these regions, because of its smoothness constraint. For the
 621 actual registration we use the software framework statismo developed at the
 622 Computer Science Department of the University of Basel. It is a framework
 623 for PCA based statistical models. These are used to describe the variability of
 624 an object within a population, learned from a set of training samples. We use
 625 it to generate a statistical model from the floating mesh. Furthermore we use
 626 the software package gpfitting for the actual fitting. We generate a infinite row
 627 of faces from the statistical model using gaussian processes and then sample
 628 out a fixed number. Then the faces are left. Carry on.

629 4.5 Prior Model

630 what to say here? describe programme?

631 4.6 Posterior Model

632 what to say here? describe programme?

633 4.7 Fitting

634 4.8 Optimizing the Loss Function

635 4.9 Varying the Variances